

Uber ride fare prediction using regression algorithms

Importing Libraries

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
```

```
In [2]: import warnings
```

```
In [3]: warnings.filterwarnings("ignore")
```

```
In [4]: data = pd.read_csv("/Users/shreyas/Downloads/RAW/codes/uber.csv")
```

Creating Data Copy

```
In [5]: df = data.copy()
```

```
In [6]: df.head()
```

```
Out[6]:
```

	Unnamed: 0	key	fare_amount	pickup_datetime	pickup_longitude	pickup_latitude
0	24238194	2015-05-07 19:52:06.0000003	7.5	2015-05-07 19:52:06 UTC	-73.999817	40.73835
1	27835199	2009-07-17 20:04:56.0000002	7.7	2009-07-17 20:04:56 UTC	-73.994355	40.72822
2	44984355	2009-08-24 21:45:00.00000061	12.9	2009-08-24 21:45:00 UTC	-74.005043	40.74077
3	25894730	2009-06-26 08:22:21.0000001	5.3	2009-06-26 08:22:21 UTC	-73.976124	40.79084
4	17610152	2014-08-28 17:47:00.000000188	16.0	2014-08-28 17:47:00 UTC	-73.925023	40.74408

In [7]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200000 entries, 0 to 199999
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Unnamed: 0            200000 non-null  int64
1   key                   200000 non-null  object
2   fare_amount           200000 non-null  float64
3   pickup_datetime       200000 non-null  object
4   pickup_longitude      200000 non-null  float64
5   pickup_latitude       200000 non-null  float64
6   dropoff_longitude     199999 non-null  float64
7   dropoff_latitude      199999 non-null  float64
8   passenger_count       200000 non-null  int64
dtypes: float64(5), int64(2), object(2)
memory usage: 13.7+ MB
```

In [8]: df["pickup_datetime"] = pd.to_datetime(df["pickup_datetime"])

In [9]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200000 entries, 0 to 199999
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Unnamed: 0            200000 non-null  int64
1   key                   200000 non-null  object
2   fare_amount           200000 non-null  float64
3   pickup_datetime       200000 non-null  datetime64[ns, UTC]
4   pickup_longitude      200000 non-null  float64
5   pickup_latitude       200000 non-null  float64
6   dropoff_longitude     199999 non-null  float64
7   dropoff_latitude      199999 non-null  float64
8   passenger_count       200000 non-null  int64
dtypes: datetime64[ns, UTC](1), float64(5), int64(2), object(1)
memory usage: 13.7+ MB
```

In [11]: `df.describe()`

Out[11]:

	Unnamed: 0	fare_amount	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude
count	2.000000e+05	200000.000000	200000.000000	200000.000000	199999.000000	199999.000000
mean	2.771250e+07	11.359955	-72.527638	39.935885	-72.525292	39.935885
std	1.601382e+07	9.901776	11.437787	7.720539	13.117408	7.720539
min	1.000000e+00	-52.000000	-1340.648410	-74.015515	-3356.666300	-88.015515
25%	1.382535e+07	6.000000	-73.992065	40.734796	-73.991407	40.734796
50%	2.774550e+07	8.500000	-73.981823	40.752592	-73.980093	40.752592
75%	4.155530e+07	12.500000	-73.967154	40.767158	-73.963658	40.767158
max	5.542357e+07	499.000000	57.418457	1644.421482	1153.572603	87.418457

In [12]: `df.isnull().sum()`


Out[12]:

Unnamed: 0	0
key	0
fare_amount	0
pickup_datetime	0
pickup_longitude	0
pickup_latitude	0
dropoff_longitude	1
dropoff_latitude	1
passenger_count	0
dtype:	int64

In [13]: `df.dropna(inplace=True)`

```
In [14]: plt.boxplot(df['fare_amount'])
```

```
Out[14]: {'whiskers': [<matplotlib.lines.Line2D at 0x165f8b890>,  
  <matplotlib.lines.Line2D at 0x165f3c6d0>],  
  'caps': [<matplotlib.lines.Line2D at 0x165f3d2d0>,  
  <matplotlib.lines.Line2D at 0x165f3e050>],  
  'boxes': [<matplotlib.lines.Line2D at 0x12f66a790>],  
  'medians': [<matplotlib.lines.Line2D at 0x165f3ea50>],  
  'fliers': [<matplotlib.lines.Line2D at 0x165f3f5d0>],  
  'means': []}
```



Remove Outliers

```
In [15]: q_low = df['fare_amount'].quantile(0.01)  
q_hi = df['fare_amount'].quantile(0.99)  
  
df = df[(df["fare_amount"] < q_hi) & (df["fare_amount"] > q_low)]
```

```
In [16]: df.isnull().sum()
```

```
Out[16]: Unnamed: 0      0
         key          0
         fare_amount    0
         pickup_datetime 0
         pickup_longitude 0
         pickup_latitude  0
         dropoff_longitude 0
         dropoff_latitude 0
         passenger_count  0
         dtype: int64
```

```
In [17]: from sklearn.model_selection import train_test_split
```

```
In [18]: X = df.drop("fare_amount", axis=1)
```

```
In [20]: y = df['fare_amount']
```

```
In [21]: X['pickup_datetime'] = pd.to_numeric(pd.to_datetime(X['pickup_datetime']))
```

```
In [22]: X = X.loc[ : , X.columns.str.contains('^Unnamed')]
```

```
In [23]: x_train, x_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, ran
```

```
In [24]: from sklearn.linear_model import LinearRegression
```

```
In [25]: lrmodel = LinearRegression()
```

```
In [26]: lrmodel.fit(x_train, y_train)
```

```
Out[26]: LinearRegression()
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [27]: predict = lrmodel.predict(x_test)
```

```
In [28]: from sklearn.metrics import mean_squared_error
```

```
In [29]: lrmodelrmse = np.sqrt(mean_squared_error(predict, y_test))
```

```
In [30]: print("RMSE error for the model is: ", lrmodelrmse)
```

RMSE error for the model is: 8.063863046328835

```
In [31]: from sklearn.ensemble import RandomForestRegressor
```

```
In [32]: rfrmodel = RandomForestRegressor(n_estimators=100, random_state=101)
```

```
In [33]: rfrmodel.fit(x_train,y_train)
```

```
Out[33]: RandomForestRegressor(random_state=101)
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [34]: rfrmodel_pred = rfrmodel.predict(x_test)
```

```
In [35]: rfrmodel_rmse = np.sqrt(mean_squared_error(rfrmodel_pred, y_test))
```

```
In [36]: print('RMSE value for random Forest is:', rfrmodel_rmse)
```

RMSE value for random Forest is: 9.757713738069647